

Computer Vision Assignment

Student Name: Stephen pangga

Student Numebr: 629869

Image resizing

```
In [ ]: from PIL import Image
import os

path = r"C:\Users\Stephen Pangga\Desktop\Computer_vision_v2\All_Test_Tubes/"
new_path = r"C:\Users\Stephen Pangga\Desktop\Computer_vision_v2\Resize"
resize_ratio = 0.5 # where 0.5 is half size

def resize_aspect_fit():
    dirs = os.listdir(path)
    for item in dirs:
        if item == '.jpg':
            continue
        if os.path.isfile(path+item):
            image = Image.open(path+item)
            file_path, extension = os.path.splitext(path+item)

            new_image_height = int(image.size[0] / (1/resize_ratio))
            new_image_length = int(image.size[1] / (1/resize_ratio))

            image = image.resize((new_image_height, new_image_length), Image.ANTIALIAS)
            image.save(file_path + "_small" + extension, 'JPEG', quality=90)

resize_aspect_fit()
```

Data Augmentaion

```
In [ ]: # pip install imageio
# pip install imgaug
# matplotlib inline
import imageio
import imgaug as ia
import imgaug.augmenters as iaa
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import matplotlib
import os
import cv2

# C:\Users\Stephen Pangga\Desktop\Computer_vision_v2\Resize tubes Organize Backup
generalpath = "/Users/Stephen Pangga/Desktop/Computer_vision_v2/Resize tubes Organize Backup"
inputpath = os.path.join(generalpath)
outputpath = os.path.join(generalpath, 'TT')

imageCountOfLetter = 0

for subdir, dirs, files in os.walk(inputpath):
    for filename in files:
```

```

if filename.endswith('.jpg'):

    imgpath = os.path.join(subdir, filename)
    filename_stripped = filename.strip('.jpg')
    outputfilepath = os.path.join(subdir, filename_stripped).replace('v1',

    print(imgpath)
    print(filename_stripped)
    print(outputfilepath)

    # 1 Default img
    defaultimg = imageio.imread(imgpath)
    img1 = cv2.resize(defaultimg, (220, 380))
    img = cv2.cvtColor(img1, cv2.COLOR_BGR2RGB)
    cv2.imwrite(outputfilepath + ".jpg", img)

    # 2 Rotate
    rotate=iaa.Affine(rotate=(-30, 30))
    rotated_image=rotate.augment_image(img)
    cv2.imwrite(outputfilepath + "_rotated.jpg", rotated_image)

    # 3 Rotate
    rotate=iaa.Affine(rotate=(30, -30))
    rotated_image=rotate.augment_image(img)
    cv2.imwrite(outputfilepath + "_rotated_2.jpg", rotated_image)

    rotate=iaa.Affine(rotate=(-50, 50))
    rotated_image=rotate.augment_image(img)
    cv2.imwrite(outputfilepath + "_rotated_3.jpg", rotated_image)

    rotate=iaa.Affine(rotate=(50, -50))
    rotated_image=rotate.augment_image(img)
    cv2.imwrite(outputfilepath + "_rotated_4.jpg", rotated_image)

    # 3 Gauss noise
    gaussian_noise=iaa.AdditiveGaussianNoise(10,20)
    noise_image=gaussian_noise.augment_image(img)
    cv2.imwrite(outputfilepath + "_gauss_noise.jpg", noise_image)

    gaussian_noise=iaa.AdditiveGaussianNoise(20,10)
    noise_image=gaussian_noise.augment_image(img)
    cv2.imwrite(outputfilepath + "_gauss_noise_1.jpg", noise_image)

    # 4 Cropping
    crop = iaa.Crop(percent=(0, 0.3)) # crop image
    crop_image = crop.augment_image(img)
    cv2.imwrite(outputfilepath + "_crop.jpg", crop_image)

    # 5 Shearing
    shear = iaa.Affine(shear=(5))
    shear_image = shear.augment_image(img)
    cv2.imwrite(outputfilepath + "_shear.jpg", shear_image)

    # 6 Flipping horizontally
    flip_hr = iaa.Fliplr(p=1.0)
    flip_hr_image = flip_hr.augment_image(img)
    cv2.imwrite(outputfilepath + "_flip_hr.jpg", flip_hr_image)

    # 7 Flip vertically
    flip_vr = iaa.Flipud(p=1.0)
    flip_vr_image= flip_vr.augment_image(img)
    cv2.imwrite(outputfilepath + "_flip_vr.jpg", flip_vr_image)

```

```

# 8 Change brightness
contrast = iaa.GammaContrast(gamma=2.0)
contrast_image = contrast.augment_image(img)
cv2.imwrite(outputfilepath + "_contrast.jpg", contrast_image)

# 9 Scaling image
scale_im = iaa.Affine(scale={"x": (1.5, 1.0), "y": (1.5, 1.0)})
scale_image = scale_im.augment_image(img)
cv2.imwrite(outputfilepath + "_scaled.jpg", scale_image)

# 10 Blur image
blurred_img = cv2.blur(img, (5, 5))
cv2.imwrite(outputfilepath + "_blurred.jpg", blurred_img)

blurred_img = cv2.blur(img, (3, 3))
cv2.imwrite(outputfilepath + "_blurred3.jpg", blurred_img)

```

Data Division Train - Validation - Test

```

In [ ]: %pip install split-folders
import os
import splitfolders
input_folder = os.path.dirname("Resize tubes Organize 220x380/")

output = "output_path_no_Plain_220x380v3"

splitfolders.ratio(input_folder, output=output, seed=42, ratio=(.7, .2, .1), group_

# splitfolders.fixed(input_folder, output=output, seed=1337, fixed=(70,10,5), overs

```

Model Training Preparation

```

In [ ]: # %pip install keras
# %pip install tensorflow
# %pip install image_dataset_loader

import keras
import os
from keras.models import *
from keras.layers import *
from keras.datasets import cifar10
from keras.optimizers import *
from keras.preprocessing.image import ImageDataGenerator
from matplotlib import pyplot as plt
from keras.utils import *
from keras.applications.vgg16 import VGG16
from image_dataset_loader import load

# train_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/train/")
# validation_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/val/")
# test_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/test/")

train_data_dir = os.path.dirname("output_path/train/")
validation_data_dir = os.path.dirname("output_path/val/")
test_data_dir = os.path.dirname("output_path/test/")

```

```

In [ ]: img_width, img_height = 224, 224

```

```

batch_size = 28

datagenerate_train = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

train_generator = datagenerate_train.flow_from_directory(train_data_dir,
                                                         target_size=(img_width,img_height),
                                                         batch_size=batch_size,
                                                         #subset="training",
                                                         class_mode='categorical')

datagenerate_validation = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

validation_generator = datagenerate_validation.flow_from_directory(validation_data_dir,
                                                                    target_size=(img_width,img_height),
                                                                    batch_size=batch_size,
                                                                    #subset="validation",
                                                                    class_mode='categorical')

print(train_generator)
print(validation_generator)
# print(test_generator)

```

Found 816 images belonging to 7 classes.
 Found 101 images belonging to 7 classes.
 <keras.preprocessing.image.DirectoryIterator object at 0x000001650A0A5850>
 <keras.preprocessing.image.DirectoryIterator object at 0x0000016504588D90>

The Model

```

In [ ]: def define_VGGmodel():

    model = VGG16(include_top=False, input_shape=(img_width, img_height, 3))

    for layer in model.layers:
        layer.trainable = False

    flat1 = Flatten()(model.layers[-1].output)

    class1 = Dense(128, activation='relu', kernel_initializer='he_uniform')(flat1)

    output = Dense(7, activation='softmax')(class1)

    model = Model(inputs=model.inputs, outputs=output)

    #opt = SGD(Lr=0.001, momentum=0.9)

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['categorical_accuracy'])

    return model

model = define_VGGmodel()
model.summary()

```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_2 (Dense)	(None, 128)	3211392
dense_3 (Dense)	(None, 7)	903
Total params: 17,926,983		
Trainable params: 3,212,295		
Non-trainable params: 14,714,688		

```
In [ ]: from keras import callbacks
        earlystopping = callbacks.EarlyStopping(monitor="val_loss",
                                                mode="min", patience=5,
                                                restore_best_weights=True)
```

```
In [ ]: history = model.fit(train_generator, steps_per_epoch=len(train_generator), validation_generator=validation_generator, validation_steps=len(validation_generator), epochs=20, verbose=1, shuffle=False, ...)
```

```

Epoch 1/20
30/30 [=====] - 87s 3s/step - loss: 2.8826 - categorical_
accuracy: 0.2132 - val_loss: 1.7610 - val_categorical_accuracy: 0.2772
Epoch 2/20
30/30 [=====] - 85s 3s/step - loss: 1.5195 - categorical_
accuracy: 0.4118 - val_loss: 1.4664 - val_categorical_accuracy: 0.3762
Epoch 3/20
30/30 [=====] - 82s 3s/step - loss: 1.1327 - categorical_
accuracy: 0.6998 - val_loss: 1.1365 - val_categorical_accuracy: 0.6634
Epoch 4/20
30/30 [=====] - 80s 3s/step - loss: 0.8550 - categorical_
accuracy: 0.8297 - val_loss: 0.9714 - val_categorical_accuracy: 0.7426
Epoch 5/20
30/30 [=====] - 82s 3s/step - loss: 0.6739 - categorical_
accuracy: 0.8725 - val_loss: 0.8603 - val_categorical_accuracy: 0.7327
Epoch 6/20
30/30 [=====] - 80s 3s/step - loss: 0.5375 - categorical_
accuracy: 0.8909 - val_loss: 0.7018 - val_categorical_accuracy: 0.8713
Epoch 7/20
30/30 [=====] - 79s 3s/step - loss: 0.4243 - categorical_
accuracy: 0.9277 - val_loss: 0.6256 - val_categorical_accuracy: 0.8416
Epoch 8/20
30/30 [=====] - 80s 3s/step - loss: 0.3339 - categorical_
accuracy: 0.9522 - val_loss: 0.5824 - val_categorical_accuracy: 0.8416
Epoch 9/20
30/30 [=====] - 80s 3s/step - loss: 0.2777 - categorical_
accuracy: 0.9669 - val_loss: 0.4947 - val_categorical_accuracy: 0.8713
Epoch 10/20
30/30 [=====] - 81s 3s/step - loss: 0.2219 - categorical_
accuracy: 0.9816 - val_loss: 0.4690 - val_categorical_accuracy: 0.9109
Epoch 11/20
30/30 [=====] - 80s 3s/step - loss: 0.1745 - categorical_
accuracy: 0.9890 - val_loss: 0.4478 - val_categorical_accuracy: 0.8515
Epoch 12/20
30/30 [=====] - 80s 3s/step - loss: 0.1473 - categorical_
accuracy: 0.9975 - val_loss: 0.3961 - val_categorical_accuracy: 0.8911
Epoch 13/20
30/30 [=====] - 81s 3s/step - loss: 0.1226 - categorical_
accuracy: 0.9951 - val_loss: 0.3898 - val_categorical_accuracy: 0.9109
Epoch 14/20
30/30 [=====] - 80s 3s/step - loss: 0.1050 - categorical_
accuracy: 0.9975 - val_loss: 0.3543 - val_categorical_accuracy: 0.9109
Epoch 15/20
30/30 [=====] - 81s 3s/step - loss: 0.0876 - categorical_
accuracy: 0.9988 - val_loss: 0.3392 - val_categorical_accuracy: 0.8911
Epoch 16/20
30/30 [=====] - 81s 3s/step - loss: 0.0740 - categorical_
accuracy: 0.9988 - val_loss: 0.3395 - val_categorical_accuracy: 0.8812
Epoch 17/20
30/30 [=====] - 87s 3s/step - loss: 0.0651 - categorical_
accuracy: 1.0000 - val_loss: 0.3187 - val_categorical_accuracy: 0.9208
Epoch 18/20
30/30 [=====] - 81s 3s/step - loss: 0.0617 - categorical_
accuracy: 1.0000 - val_loss: 0.3216 - val_categorical_accuracy: 0.8812
Epoch 19/20
30/30 [=====] - 80s 3s/step - loss: 0.0498 - categorical_
accuracy: 1.0000 - val_loss: 0.3154 - val_categorical_accuracy: 0.8812
Epoch 20/20
30/30 [=====] - 81s 3s/step - loss: 0.0437 - categorical_
accuracy: 1.0000 - val_loss: 0.3250 - val_categorical_accuracy: 0.9010


```

```
In [ ]: print(history.history)
```


```
{'loss': [2.882565975189209, 1.5194791555404663, 1.1327366828918457, 0.85503476858
13904, 0.6738737225532532, 0.5374535322189331, 0.42425525188446045, 0.333894014358
5205, 0.27769342064857483, 0.22190436720848083, 0.17446686327457428, 0.14725792407
989502, 0.1226036474108696, 0.1049560010433197, 0.087623231112957, 0.0739550739526
7487, 0.06513470411300659, 0.0616726316511631, 0.0498480498790741, 0.0436756089329
71954], 'categorical_accuracy': [0.2132352888584137, 0.4117647111415863, 0.6997548
937797546, 0.8296568393707275, 0.8725489974021912, 0.8909313678741455, 0.927696049
2134094, 0.9522058963775635, 0.966911792755127, 0.9816176295280457, 0.988970577716
8274, 0.9975489974021912, 0.9950980544090271, 0.9975489974021912, 0.99877452850341
8, 0.998774528503418, 1.0, 1.0, 1.0, 1.0], 'val_loss': [1.7610087394714355, 1.4664
227962493896, 1.136460542678833, 0.9714139699935913, 0.8603437542915344, 0.7017602
324485779, 0.6255902647972107, 0.5824043154716492, 0.49471667408943176, 0.46903523
802757263, 0.44778385758399963, 0.3960905969142914, 0.3898434638977051, 0.35432577
13317871, 0.33916419744491577, 0.3394950330257416, 0.31873181462287903, 0.32155048
847198486, 0.31542450189590454, 0.3250230550765991], 'val_categorical_accuracy':
[0.2772277295589447, 0.3762376308441162, 0.6633663177490234, 0.7425742745399475,
0.7326732873916626, 0.8712871074676514, 0.8415841460227966, 0.8415841460227966, 0.
8712871074676514, 0.9108911156654358, 0.8514851331710815, 0.8910890817642212, 0.91
08911156654358, 0.9108911156654358, 0.8910890817642212, 0.8811880946159363, 0.9207
921028137207, 0.8811880946159363, 0.8811880946159363, 0.9009901285171509]}
```

Diagnostic plot

```
In [ ]: plt.plot(history.history['categorical_accuracy'])
plt.plot(history.history['val_categorical_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val_loss'], loc='upper left')
plt.show()
```



```
In [ ]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Loss', 'Val_loss'], loc='upper right')
plt.show()
```



```
In [ ]: print('Validation image - evaluation')
_, acc = model.evaluate(validation_generator, steps=len(validation_generator), verbose=1)
print('accuracy score: ', acc*100 )
```

Validation image - evaluation
 4/4 [=====] - 10s 2s/step - loss: 0.3250 - categorical_accuracy: 0.9010
 accuracy score: 90.09901285171509

```
In [ ]: # print('Test image - evaluation')
# _, acc = model.evaluate(test_generator, steps=len(test_generator), verbose=1)
# print('accuracy score: ', acc*100 )
```

Testing

```
In [ ]: import numpy as np
from sklearn.metrics import confusion_matrix
from sklearn import metrics
# test_data_dir = os.path.dirname("output_path_no_PLain_220x380v1/train-test/")

datagenerate_test = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

test_generator = datagenerate_test.flow_from_directory(test_data_dir,
                                                        target_size=(img_width,img_height),
                                                        batch_size=batch_size,
                                                        #subset="validation",
                                                        class_mode='categorical')
```

Found 103 images belonging to 7 classes.

```
In [ ]: print('Test image - evaluation')
_, acc = model.evaluate(test_generator, steps=len(test_generator), verbose=1)
print('accuracy score: ', acc*100 )
```

Test image - evaluation
 4/4 [=====] - 9s 2s/step - loss: 0.2666 - categorical_accuracy: 0.9320
 accuracy score: 93.20388436317444

```
In [ ]: Y_pred = model.predict(test_generator)
y_pred = np.argmax(Y_pred, axis=1)
print('accuracy: ', metrics.accuracy_score(y_pred,test_generator.classes))
```

4/4 [=====] - 9s 2s/step
 accuracy: 0.1553398058252427



```
In [ ]: print('Confusion Matrix')
print(confusion_matrix(y_pred, test_generator.classes))
```

```
Confusion Matrix
[[2 2 2 3 1 3 2]
 [4 3 4 2 2 0 2]
 [1 4 0 1 2 3 0]
 [1 1 2 0 1 2 1]
 [1 3 3 1 3 4 3]
 [3 4 1 1 5 4 2]
 [2 1 0 1 4 2 4]]
```

```
In [ ]: import seaborn as sns
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

cm = confusion_matrix(y_pred, test_generator.classes)
cmd = ConfusionMatrixDisplay(cm, display_labels=[ 'BD-Vacutainer', 'EDTA(K2)', 'Monovette', 'SGS', 'Vacuette', 'Vacutest', 'Venosafe'])
fig, ax = plt.subplots(figsize=(10,10))
plt.grid(False)
cmd.plot(ax=ax)
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1651675dd60>
```



```
In [ ]: from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

print('Accuracy Score :', accuracy_score(y_pred, test_generator.classes))
print('Report : ')
target_names = ['BD-Vacutainer', 'EDTA(K2)', 'Monovette', 'SGS', 'Vacuette', 'Vacutest', 'Venosafe']
print(classification_report(y_pred, test_generator.classes, target_names=target_names))
```

Accuracy Score : 0.1553398058252427

Report :

	precision	recall	f1-score	support
BD-Vacutainer	0.14	0.13	0.14	15
EDTA(K2)	0.17	0.18	0.17	17
Monovette	0.00	0.00	0.00	11
SGS	0.00	0.00	0.00	8
Vacurette	0.17	0.17	0.17	18
Vacutest	0.22	0.20	0.21	20
Venosafe	0.29	0.29	0.29	14
accuracy			0.16	103
macro avg	0.14	0.14	0.14	103
weighted avg	0.16	0.16	0.16	103

The Accuracy score is 15.53%, this could be due to the image. One of the plausible issues i can think of that cause such a low accuracy is the fact that maybe the augmented image type has not been since by the model.