

Computer vision assignment

Student name: Stephen Pangga

Student number: 629860

Imports

```
In [ ]: # %pip install keras
# %pip install tensorflow
# %pip install image_dataset_loader

import keras
import os
from keras.models import *
from keras.layers import *
from keras.datasets import cifar10
from keras.optimizers import *
from keras.preprocessing.image import ImageDataGenerator
from matplotlib import pyplot as plt
from keras.utils import *
from keras.applications.vgg16 import VGG16
from keras.applications.vgg19 import VGG19

from image_dataset_loader import load

# train_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/train/")
# validation_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/val/")
# test_data_dir = os.path.dirname("output_path_no_Plain_220x380v3/test/")

train_data_dir = os.path.dirname("output_path/train/")
validation_data_dir = os.path.dirname("output_path/val/")
test_data_dir = os.path.dirname("output_path/test/")

In [ ]: img_width, img_height = 224, 224
batch_size = 16

datagenerate_train = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

train_generator = datagenerate_train.flow_from_directory(train_data_dir,
                                                         target_size=(img_width, img_height),
                                                         batch_size=batch_size,
                                                         #subset="training",
                                                         class_mode='categorical',
                                                         shuffle = True)

datagenerate_validation = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

validation_generator = datagenerate_validation.flow_from_directory(validation_data_dir,
                                                                    target_size=(img_width, img_height),
                                                                    batch_size=batch_size,
                                                                    #subset="validation",
                                                                    class_mode='categorical',
                                                                    shuffle = True)
```

```
print(train_generator)
print(validation_generator)
# print(test_generator)
```

Found 816 images belonging to 7 classes.

Found 101 images belonging to 7 classes.

<keras.preprocessing.image.DirectoryIterator object at 0x00000246A0968160>

<keras.preprocessing.image.DirectoryIterator object at 0x00000246F90F85B0>

The Model

```
In [ ]: def define_VGGmodel():

    model = VGG16(include_top=False, input_shape=(img_width, img_height, 3))

    for layer in model.layers:
        layer.trainable = False

    flat1 = Flatten()(model.layers[-1].output)

    class1 = Dense(128, activation='relu', kernel_initializer='he_uniform')(flat1)

    output = Dense(7, activation='softmax')(class1)

    model = Model(inputs=model.inputs, outputs=output)

    # opt = SGD(lr=0.001, momentum=0.9)

    model.compile(optimizer="Adam", loss='categorical_crossentropy', metrics=['categorical_accuracy'])

    return model

model = define_VGGmodel()
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 128)	3211392
dense_1 (Dense)	(None, 7)	903
=====		
Total params: 17,926,983		
Trainable params: 3,212,295		
Non-trainable params: 14,714,688		

```
In [ ]: from keras import callbacks
        earlystopping = callbacks.EarlyStopping(monitor="val_loss",
                                                mode="min", patience=5,
                                                restore_best_weights=True)
```

```
In [ ]: history = model.fit(train_generator, steps_per_epoch=len(train_generator), validation_generator=validation_generator, validation_steps=len(validation_generator), epochs=25, verbose=1, shuffle=False,
```

Epoch 1/25
51/51 [=====] - 83s 2s/step - loss: 2.0130 - categorical_accuracy: 0.3811 - val_loss: 1.2446 - val_categorical_accuracy: 0.5545
Epoch 2/25
51/51 [=====] - 82s 2s/step - loss: 0.8218 - categorical_accuracy: 0.7414 - val_loss: 0.8679 - val_categorical_accuracy: 0.6238
Epoch 3/25
51/51 [=====] - 75s 1s/step - loss: 0.4796 - categorical_accuracy: 0.8934 - val_loss: 0.5692 - val_categorical_accuracy: 0.8317
Epoch 4/25
51/51 [=====] - 82s 2s/step - loss: 0.2860 - categorical_accuracy: 0.9559 - val_loss: 0.5903 - val_categorical_accuracy: 0.7822
Epoch 5/25
51/51 [=====] - 81s 2s/step - loss: 0.2206 - categorical_accuracy: 0.9706 - val_loss: 0.4360 - val_categorical_accuracy: 0.8911
Epoch 6/25
51/51 [=====] - 80s 2s/step - loss: 0.1335 - categorical_accuracy: 0.9877 - val_loss: 0.3606 - val_categorical_accuracy: 0.9010
Epoch 7/25
51/51 [=====] - 80s 2s/step - loss: 0.0900 - categorical_accuracy: 0.9975 - val_loss: 0.4320 - val_categorical_accuracy: 0.8515
Epoch 8/25
51/51 [=====] - 81s 2s/step - loss: 0.0849 - categorical_accuracy: 0.9951 - val_loss: 0.3602 - val_categorical_accuracy: 0.8812
Epoch 9/25
51/51 [=====] - 83s 2s/step - loss: 0.0753 - categorical_accuracy: 0.9963 - val_loss: 0.2816 - val_categorical_accuracy: 0.9109
Epoch 10/25
51/51 [=====] - 79s 2s/step - loss: 0.0401 - categorical_accuracy: 1.0000 - val_loss: 0.2737 - val_categorical_accuracy: 0.9010
Epoch 11/25
51/51 [=====] - 80s 2s/step - loss: 0.0309 - categorical_accuracy: 1.0000 - val_loss: 0.2734 - val_categorical_accuracy: 0.8911
Epoch 12/25
51/51 [=====] - 80s 2s/step - loss: 0.0267 - categorical_accuracy: 1.0000 - val_loss: 0.2580 - val_categorical_accuracy: 0.9109
Epoch 13/25
51/51 [=====] - 81s 2s/step - loss: 0.0208 - categorical_accuracy: 1.0000 - val_loss: 0.2787 - val_categorical_accuracy: 0.9010
Epoch 14/25
51/51 [=====] - 81s 2s/step - loss: 0.0179 - categorical_accuracy: 1.0000 - val_loss: 0.2540 - val_categorical_accuracy: 0.8911
Epoch 15/25
51/51 [=====] - 81s 2s/step - loss: 0.0154 - categorical_accuracy: 1.0000 - val_loss: 0.2406 - val_categorical_accuracy: 0.9109
Epoch 16/25
51/51 [=====] - 81s 2s/step - loss: 0.0132 - categorical_accuracy: 1.0000 - val_loss: 0.2253 - val_categorical_accuracy: 0.9208
Epoch 17/25
51/51 [=====] - 80s 2s/step - loss: 0.0124 - categorical_accuracy: 1.0000 - val_loss: 0.2253 - val_categorical_accuracy: 0.9208
Epoch 18/25
51/51 [=====] - 80s 2s/step - loss: 0.0115 - categorical_accuracy: 1.0000 - val_loss: 0.2324 - val_categorical_accuracy: 0.9010
Epoch 19/25
51/51 [=====] - 85s 2s/step - loss: 0.0092 - categorical_accuracy: 1.0000 - val_loss: 0.2204 - val_categorical_accuracy: 0.9109
Epoch 20/25
51/51 [=====] - 89s 2s/step - loss: 0.0083 - categorical_accuracy: 1.0000 - val_loss: 0.2336 - val_categorical_accuracy: 0.8911
Epoch 21/25
51/51 [=====] - 81s 2s/step - loss: 0.0072 - categorical_accuracy: 1.0000 - val_loss: 0.2134 - val_categorical_accuracy: 0.9010
Epoch 22/25

```

51/51 [=====] - 78s 2s/step - loss: 0.0067 - categorical_
accuracy: 1.0000 - val_loss: 0.2289 - val_categorical_accuracy: 0.9010
Epoch 23/25
51/51 [=====] - 77s 2s/step - loss: 0.0061 - categorical_
accuracy: 1.0000 - val_loss: 0.2094 - val_categorical_accuracy: 0.9109
Epoch 24/25
51/51 [=====] - 78s 2s/step - loss: 0.0055 - categorical_
accuracy: 1.0000 - val_loss: 0.2110 - val_categorical_accuracy: 0.9109
Epoch 25/25
51/51 [=====] - 77s 2s/step - loss: 0.0051 - categorical_
accuracy: 1.0000 - val_loss: 0.2282 - val_categorical_accuracy: 0.8911

```

```
In [ ]: print(history.history)
```


```

{'loss': [2.012996196746826, 0.8217641711235046, 0.47962120175361633, 0.2860125303
2684326, 0.2206156849861145, 0.13351689279079437, 0.08996354043483734, 0.084902554
7504425, 0.07527531683444977, 0.04009522125124931, 0.03093503974378109, 0.02673844
8068499565, 0.02081669308245182, 0.017879987135529518, 0.015433376654982567, 0.013
170113787055016, 0.012445982545614243, 0.011478454805910587, 0.009217247366905212,
0.0083229411393404, 0.007217416074126959, 0.006735136266797781, 0.0060526682063937
19, 0.005545631516724825, 0.005145586561411619], 'categorical_accuracy': [0.381127
4468898773, 0.7414215803146362, 0.8933823704719543, 0.9558823704719543, 0.97058820
7244873, 0.9877451062202454, 0.9975489974021912, 0.9950980544090271, 0.99632352590
56091, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0,
1.0], 'val_loss': [1.2445828914642334, 0.8678803443908691, 0.5691992044448853, 0.5
90302586555481, 0.4359528124332428, 0.3605560064315796, 0.43199849128723145, 0.360
18794775009155, 0.28156766295433044, 0.2736772298812866, 0.2734403908252716, 0.257
9731047153473, 0.2786790132522583, 0.2539757192134857, 0.2405688762664795, 0.22528
916597366333, 0.22528526186943054, 0.2323930710554123, 0.220425084233284, 0.233551
0402917862, 0.21336442232131958, 0.22891883552074432, 0.2094143182039261, 0.210956
409573555, 0.22817252576351166], 'val_categorical_accuracy': [0.5544554591178894,
0.6237623691558838, 0.8316831588745117, 0.7821782231330872, 0.8910890817642212, 0.
9009901285171509, 0.8514851331710815, 0.8811880946159363, 0.9108911156654358, 0.90
09901285171509, 0.8910890817642212, 0.9108911156654358, 0.9009901285171509, 0.8910
890817642212, 0.9108911156654358, 0.9207921028137207, 0.9207921028137207, 0.900990
1285171509, 0.9108911156654358, 0.8910890817642212, 0.9009901285171509, 0.90099012
85171509, 0.9108911156654358, 0.9108911156654358, 0.8910890817642212]}


```

Diagnostic plot

```
In [ ]: plt.plot(history.history['categorical_accuracy'])
plt.plot(history.history['val_categorical_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val_loss'], loc='upper left')
plt.show()
```



```
In [ ]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Loss', 'Val_loss'], loc='upper right')
plt.show()
```



```
In [ ]: print('Validation image - evaluation')
_, acc = model.evaluate(validation_generator, steps=len(validation_generator), verbose=0)
print('accuracy score: ', acc*100)
```

Validation image - evaluation

7/7 [=====] - 9s 1s/step - loss: 0.2282 - categorical_accuracy: 0.8911

accuracy score: 89.10890817642212

```
In [ ]: model.save('VGG16-accu-score-92%')
```

```
#model=Load_model('vgg15-softmax-classifier_train-91.54_val-72.06_test-83.82%')
```

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 13). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: VGG16-accu-score-92%\assets

INFO:tensorflow:Assets written to: VGG16-accu-score-92%\assets

Testing

```
In [ ]: import numpy as np
from sklearn.metrics import confusion_matrix
from sklearn import metrics
# test_data_dir = os.path.dirname("output_path_no_Plain_220x380v1/train-test/")

datagenerate_test = ImageDataGenerator(rescale=1.0/255, samplewise_center=True)

test_generator = datagenerate_test.flow_from_directory(test_data_dir,
                                                        target_size=(img_width,img_height),
                                                        batch_size=batch_size,
                                                        #subset="validation",
                                                        class_mode='categorical',
                                                        shuffle = False)
```

Found 103 images belonging to 7 classes.

```
In [ ]: #model=Load_model('VGG16-accu-score-92%')
```

```
In [ ]: print('Test image - evaluation')
_, acc = model.evaluate(test_generator, steps=len(test_generator), verbose=1)
print('accuracy score: ', acc*100 )

Test image - evaluation
7/7 [=====] - 10s 1s/step - loss: 0.1904 - categorical_ac
curacy: 0.9126
accuracy score: 91.26213788986206
```

```
In [ ]: Y_pred = model.predict(test_generator)
y_pred = np.argmax(Y_pred, axis=1)
print('accuracy: ', metrics.accuracy_score(y_pred,test_generator.classes))

7/7 [=====] - 9s 1s/step
accuracy: 0.912621359223301
```


```
In [ ]: print('Confusion Matrix')
print(confusion_matrix(y_pred, test_generator.classes))

Confusion Matrix
[[13  0  0  1  0  0  0]
 [ 0 17  0  0  1  0  0]
 [ 0  0 12  0  0  0  0]
 [ 0  0  0  8  0  0  0]
 [ 1  0  0  0 16  0  0]
 [ 0  1  0  0  1 17  3]
 [ 0  0  0  0  0  1 11]]
```

```
In [ ]: import seaborn as sns
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

cm = confusion_matrix(y_pred, test_generator.classes)
cmd = ConfusionMatrixDisplay(cm, display_labels=[ 'BD-Vacutainer', 'EDTA(K2)', 'Moi
fig, ax = plt.subplots(figsize=(10,10))
plt.grid(False)
cmd.plot(ax=ax)
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x246b68c9880>
```



```
In [ ]: from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

print('Accuracy Score :', accuracy_score(y_pred, test_generator.classes))
print ('Report : ')
target_names = ['BD-Vacutainer', 'EDTA(K2)', 'Monovette', 'SGS', 'Vacuette' , 'Vacutest', 'Venosafe']
print (classification_report(y_pred, test_generator.classes, target_names=target_names))
```

Accuracy Score : 0.912621359223301

Report :

	precision	recall	f1-score	support
BD-Vacutainer	0.93	0.93	0.93	14
EDTA(K2)	0.94	0.94	0.94	18
Monovette	1.00	1.00	1.00	12
SGS	0.89	1.00	0.94	8
Vacuette	0.89	0.94	0.91	17
Vacutest	0.94	0.77	0.85	22
Venosafe	0.79	0.92	0.85	12
accuracy			0.91	103
macro avg	0.91	0.93	0.92	103
weighted avg	0.92	0.91	0.91	103


The first accuracy score was 15.53%, this could be due to the image. One of the plausible issues i can think of that cause such a low accuracy is the fact that maybe tht augmented image has not been since by the model.

Update:

There was an issue where the accuracy was giving 15 percent accuracy, that has been now solved. The issue was when `model.predict()` the train folder it organized the result of the folder and doesn't work with the randomize order of `test_generator.classes`. To prevent this from happening add "shuffle= False" in the parameter of your `imageDataGenerator`.

Test an image

```
In [ ]: from skimage.transform import resize
Image = plt.imread("output_path/test/BD-Vacutainer/BD-Vacutainer1_small_contrast.")
RImage = resize(Image, (img_width, img_height, 3))
img = plt.imshow(RImage)
```



```
In [ ]: predictions = model.predict(np.array( [RImage] ))
predictions

y_pred = np.argmax(predictions, axis=1)

print(y_pred)
# should give 0 as it is vacutainer

1/1 [=====] - 0s 192ms/step
[0]
```